


# Hyperbolic Restricted Boltzmann Machines for Hierarchical Learning

Kylie Hefner

Chapman University: Computational and Data Sciences  
CS 595: Computational Science Seminars

# Overview

- Problem:
  - Traditional RBMs use Euclidean geometry, which distorts hierarchical/tree-like data.
  - Key Gap: Can hyperbolic geometry improve RBMs for hierarchical structures?
- Solution:
  - A **Hyperbolic RBM** with adjusted energy function and sampling.
- Impact:
  - Better unsupervised learning for NLP, biology, or knowledge graphs.




# Literature Review

- Masaki Kobayashi
  - Researches the use of non-Euclidean geometry in Hopfield Neural Networks (complex, hyperbolic, quaternion, rotor, etc.)
- John J. Hopfield
  - Invented the Hopfield Neural Network (1982)
  - 2024 Nobel Prize in Physics
- Geoffrey Hinton
  - Invented the Boltzmann Machine (1983-1985)
  - 2024 Nobel Prize in Physics



# What is a Boltzmann Machine

- Stochastic, generative neural networks that learn patterns by minimizing an energy function
- Inspired by statistical mechanics (Boltzmann distribution)
- Why they matter:
  - Foundation for energy-based models, latent variable learning, and modern generative AI.





# Energy Computation

$$E(\mathbf{v}, \mathbf{h}) = -\mathbf{v}^T W \mathbf{h} - \mathbf{a}^T \mathbf{v} - \mathbf{b}^T \mathbf{h}$$

Measures the "goodness" of a configuration (lower energy = higher probability)

Learning involves adjusting weights to reduce energy for training data

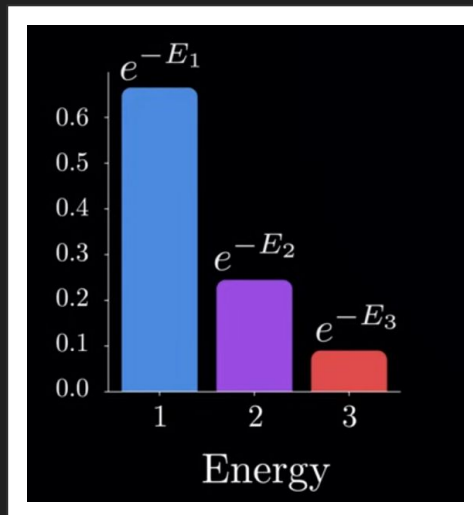


# Energy Computation

- The energy function is used to compute the **probability distribution** over all possible configurations of the units.
- Specifically, the probability of a configuration is given by the Boltzmann distribution:


$$p(\mathbf{v}, \mathbf{h}) = \frac{\exp(-E(\mathbf{v}, \mathbf{h}))}{Z}$$

( $Z$  is a normalizing function, the sum of all possible state probabilities)



# Restricted Boltzmann Machines

- Structure:
  - A simplified Boltzmann Machine with restricted architecture
  - No intra-layer connections: only visible-to-hidden links
- Training:
  - Contrastive Divergence (CD-k): Approximate gradient using Gibbs sampling (fast, but biased)



# Boltzmann Machine Simple Example

Let's say you have a dataset of movie preferences for three people. Each person rates whether they like (1) or dislike (0) two movies.

The dataset consists of the following vectors:

- $x_1 = [1, 0]$  (Person 1 likes Movie A and dislikes Movie B)
- $x_2 = [0, 1]$  (Person 2 dislikes Movie A and likes Movie B)
- $x_3 = [1, 1]$  (Person 3 likes Movie A and Movie B)


Our goal is to train a Boltzmann Machine to learn the underlying patterns in this dataset.



# Boltzmann Machine Simple Example

We create the architecture to represent this data:

- **Visible units** represent the movie preferences (movie A/B)
- **Hidden units** capture latent (unmeasured) features
- **Weights** connect units (pos/neg relationship between units)
- **Visible** and **hidden** units have biases that measure individual probability of being on or off




Initial values are chosen randomly or from domain knowledge

# Restricted Boltzmann Machine Simple Example

We create the architecture to represent this data:

- **Visible units** represent the movie preferences (movie A/B)
- **Hidden units** capture latent (unmeasured) features
- **Weights** connect units (pos/neg relationship between units)
- **Visible** and **hidden** units have biases that measure individual probability of being on or off



Initial values are chosen randomly or from domain knowledge

# Training RBMs: Steps

0. Initialize weights and biases


1. Forward Pass  
(Visible  $\rightarrow$  Hidden)

2. Reconstruct Input  
(Hidden  $\rightarrow$  Visible)

4. Update weights  
and biases

3. Forward Pass Again  
(Reconstructed Visible  
 $\rightarrow$  Hidden)

5. Repeat steps 1-4 for all data points and multiple epochs until convergence



# Training RBMs: Approximation

## Theoretical:

- Goal: Maximize the likelihood of the training data under the RBM's energy-based model
- **Problem:** Requires the partition function  $Z$ , which sums up all possible configurations of the model. Too many calculations for large models.



## Contrastive Divergence:

- **Solution:** Approximate the gradient by running Gibbs Sampling and stopping before equilibrium
- Why it works: Just a few steps of Gibbs tells us the direction of gradient needed to reduce energy.

## Step 1: Forward Pass

Sampling is used to infer hidden states given visible states

$$p(h_j = 1|\mathbf{v}) = \sigma \left( b_j + \sum_i v_i w_{ij} \right)$$

$$\sigma(x) = \frac{1}{1+e^{-x}}$$

Sampling for  $\mathbf{v} = [1, 0]$  (person 1):

**h1:**  $\sigma(-0.3 + 1 \cdot 0.5 + 0 \cdot 0.7) = \sigma(0.2) \approx 0.5498$ . Sampled  $h_1 = 1$ .

**h2:**  $\sigma(0.4 + 1 \cdot (-0.6) + 0 \cdot (-0.8)) = \sigma(-0.2) \approx 0.4502$ . Sampled  $h_2 = 0$ .

Therefore, our hidden activation is  $\mathbf{h} = [1, 0]$  for this visible state.

## Step 2: Reconstruct Input

Now, we sample visible nodes (input) using  $h = [1, 0]$  from previous step:

$$P(v_i = 1|h) = \sigma(a_i + \sum_j h_j W_{ij})$$

**v1:**  $\sigma(0.1 + 1 \cdot 0.5 + 0 \cdot (-0.6)) = \sigma(0.6) \approx 0.6457$ . Sampled  $v'_1 = 1$ .

**v2:**  $\sigma(-0.2 + 1 \cdot 0.7 + 0 \cdot (-0.8)) = \sigma(0.5) \approx 0.6225$ . Sampled  $v'_2 = 1$ .

Therefore, our reconstructed input is  $v' = [1, 1]$ .

### Step 3: Forward Pass (again)

$$p(h_j = 1|\mathbf{v}) = \sigma \left( b_j + \sum_i v_i w_{ij} \right)$$

$$\sigma(x) = \frac{1}{1+e^{-x}}$$

Sampling for  $\mathbf{v}' = [1, 1]$ :

**h1'**:  $\sigma(-0.3 + 1 \cdot 0.5 + 1 \cdot 0.7) = \sigma(0.9) \approx 0.7109$ . Sampled  $h'_1 = 1$ .

**h2'**:  $\sigma(0.4 + 1 \cdot (-0.6) + 1 \cdot (-0.8)) = \sigma(-1.0) \approx 0.2689$ . Sampled  $h'_2 = 0$ .

Therefore, our reconstructed hidden activation is  $\mathbf{h}' = [1, 0]$ .

## Step 4: Update Parameters

Goal: Adjust weights and biases to minimize difference between  $v$  and  $v'$

$$\Delta W = \epsilon (\mathbf{v}^T \mathbf{h} - \mathbf{v}'^T \mathbf{h}')$$

$$= \epsilon \left( \begin{bmatrix} 1 \\ 0 \end{bmatrix} \begin{bmatrix} 1 & 0 \end{bmatrix} - \begin{bmatrix} 1 \\ 1 \end{bmatrix} \begin{bmatrix} 1 & 0 \end{bmatrix} \right)$$

$$= 0.1 \left( \begin{bmatrix} 1 & 0 \\ 0 & 0 \end{bmatrix} - \begin{bmatrix} 1 & 0 \\ 1 & 0 \end{bmatrix} \right)$$

$$= \begin{bmatrix} 0 & 0 \\ -0.1 & 0 \end{bmatrix}$$

$$\Delta \mathbf{a} = \epsilon (\mathbf{v} - \mathbf{v}')$$

$$= 0.1 \left( \begin{bmatrix} 1 & 0 \end{bmatrix} - \begin{bmatrix} 1 & 1 \end{bmatrix} \right)$$

$$= \begin{bmatrix} 0 & -0.1 \end{bmatrix}$$

$$\Delta \mathbf{b} = \epsilon (\mathbf{h} - \mathbf{h}')$$

$$= 0.1 \left( \begin{bmatrix} 1 & 0 \end{bmatrix} - \begin{bmatrix} 1 & 0 \end{bmatrix} \right)$$

$$= \begin{bmatrix} 0 & 0 \end{bmatrix}$$



## Step 4: Update Parameters

Now, we can update the original parameters

$$W_{\text{new}} = W + \Delta W = \begin{bmatrix} 0.5 & -0.6 \\ 0.7 & -0.8 \end{bmatrix} + \begin{bmatrix} 0 & 0 \\ -0.1 & 0 \end{bmatrix} = \begin{bmatrix} 0.5 & -0.6 \\ 0.6 & -0.8 \end{bmatrix}$$

$$\mathbf{a}_{\text{new}} = \mathbf{a} + \Delta \mathbf{a} = \begin{bmatrix} 0.1 & -0.2 \end{bmatrix} + \begin{bmatrix} 0 & -0.1 \end{bmatrix} = \begin{bmatrix} 0.1 & -0.3 \end{bmatrix}$$

$$\mathbf{b}_{\text{new}} = \mathbf{b} + \Delta \mathbf{b} = \begin{bmatrix} -0.3 & 0.4 \end{bmatrix} + \begin{bmatrix} 0 & 0 \end{bmatrix} = \begin{bmatrix} -0.3 & 0.4 \end{bmatrix} \quad (\text{no change})$$

Step 5: Repeat steps 1-4 for all data points and multiple epochs until convergence


# Legacy in Modern Machine Learning

- Deep Belief Networks (DBNs):
  - Stacked RBMs → Unsupervised pre-training → Fine-tuning with backprop
  - Pioneered the "deep learning revolution" (2006–2012)
- Connections to Modern Models:
  - Energy-based models (e.g., VAEs, diffusion models)
  - Stochastic units in GANs/VAEs vs. BM's Gibbs sampling
- Why They Faded:
  - Computationally expensive vs. backpropagation-friendly architectures (CNNs, Transformers)

**Key Takeaway: Understanding Boltzmann Machines reveals the DNA of modern generative AI**


# A Primer on Hyperbolic Geometry

- A non-Euclidean geometry with constant negative curvature (like a saddle or pringle shape).
- Contrast with:
  - Euclidean (flat) space: Zero curvature.
  - Spherical geometry: Positive curvature (like a sphere).
- In hyperbolic space, distances grow exponentially as you move outward.




# Hyperbolic Models: A Tool for Hierarchical Data

- Represent data points in hyperbolic space (e.g., Poincaré ball)
- Distance between points: 
$$d(x, y) = \operatorname{arcosh} \left( 1 + 2 \frac{\|x - y\|^2}{(1 - \|x\|^2)(1 - \|y\|^2)} \right)$$
- Efficient representation of hierarchies (fewer dimensions, lower distortion).




# WordNet Hierarchy


- A lexical database where nouns are organized in a hierarchical taxonomy (e.g., "animal" → "mammal" → "dog").
- Tree-like structure with clear parent-child relationships.
- Access: `nlk.corpus.wordnet`



		living being	feline	human	gender	royalty	verb	plural
<i>man</i>	→	0.6	-0.2	0.8	0.9	-0.1	-0.9	-0.7
<i>woman</i>	→	0.7	0.3	0.8	-0.7	0.1	-0.5	-0.4
<i>king</i>	→	0.5	-0.4	0.7	0.8	0.9	-0.7	-0.6
<i>queen</i>	→	0.8	-0.1	0.8	-0.9	0.8	-0.5	-0.9
word		Word embedding						



Visualization of word embedding



# Model Design: Energy Function

## Euclidean RBM:

$$E(\mathbf{v}, \mathbf{h}) = -\mathbf{v}^T W \mathbf{h} - \mathbf{a}^T \mathbf{v} - \mathbf{b}^T \mathbf{h}$$

- $\mathbf{v}$ : Visible units.
- $\mathbf{h}$ : Hidden units.
- $W$ : Weight matrix.
- $\mathbf{a}, \mathbf{b}$ : Bias vectors.

## Hyperbolic RBM:

$$E(\mathbf{v}, \mathbf{h}) = -\beta \cdot d_{\mathbb{H}}(\mathbf{v}, W \otimes \mathbf{h})^2 - \mathbf{a}^T \otimes \mathbf{v} - \mathbf{b}^T \otimes \mathbf{h}$$

- $d_{\mathbb{H}}$ : Hyperbolic distance in the Poincaré ball.
- $\otimes$ : Lorentzian inner product (analog of dot product).
- $\beta$ : Scaling factor for curvature.



# Model Design: Sampling Methods

## Euclidean RBM:

- Gibbs Sampling: Sample  $h|v$  and  $v|h$  using logistic activation
- Markov chain typically converges quickly

$$p(h_j = 1|\mathbf{v}) = \sigma \left( b_j + \sum_i v_i w_{ij} \right)$$
$$p(v_i = 1|\mathbf{h}) = \sigma \left( a_i + \sum_j h_j w_{ij} \right)$$

## Hyperbolic RBM:

- Gibbs Sampling: Sample  $h|v$  and  $v|h$  using hyperbolic probabilities
- Markov chain may take longer to converge due to curvature

$$p(h_j = 1|\mathbf{v}) = \sigma \left( \text{logit}_{\mathbb{H}}(b_j \oplus \sum_i v_i \otimes w_{ij}) \right)$$
$$p(v_i = 1|\mathbf{h}) = \sigma \left( \text{logit}_{\mathbb{H}}(a_i \oplus \sum_j h_j \otimes w_{ij}) \right)$$

- $\oplus$ : Möbius addition.
- $\text{logit}_{\mathbb{H}}$ : Hyperbolic logit function.

# Model Design: Training

## Euclidean RBM:

- Update weights and biases using gradient ascent:

$$\Delta w_{ij} = \epsilon (\langle v_i h_j \rangle_{\text{data}} - \langle v_i h_j \rangle_{\text{model}})$$

$\epsilon$ : Learning rate.

- Key Properties:
  - Gradients are computed using standard backpropagation.
  - Optimization is stable and efficient.

## Hyperbolic RBM:

- Update weights and biases using Riemannian gradient ascent:

$$\Delta w_{ij} = \epsilon \cdot \text{proj}_{\mathbb{H}} (\langle v_i h_j \rangle_{\text{data}} - \langle v_i h_j \rangle_{\text{model}})$$

$\text{proj}_{\mathbb{H}}$ : Projection onto the hyperbolic manifold.

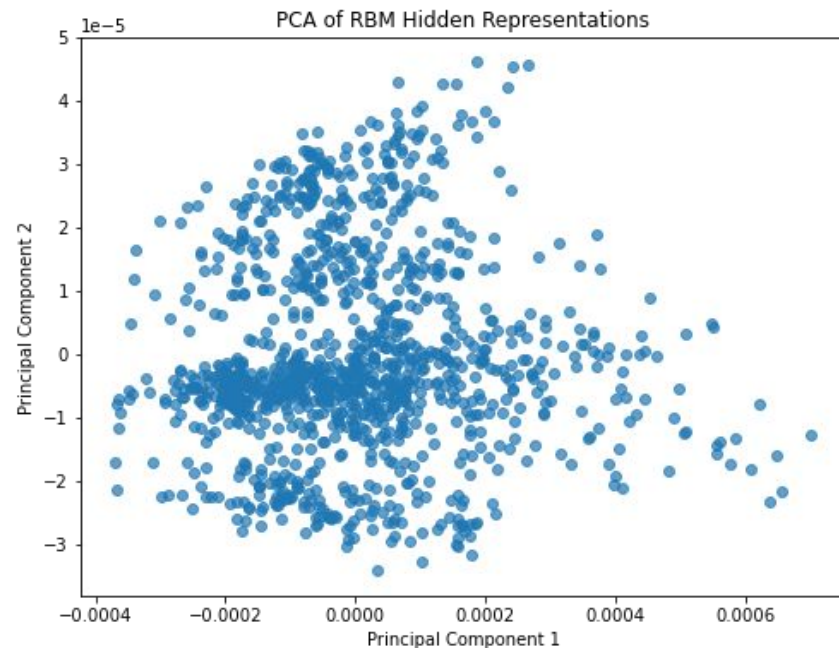
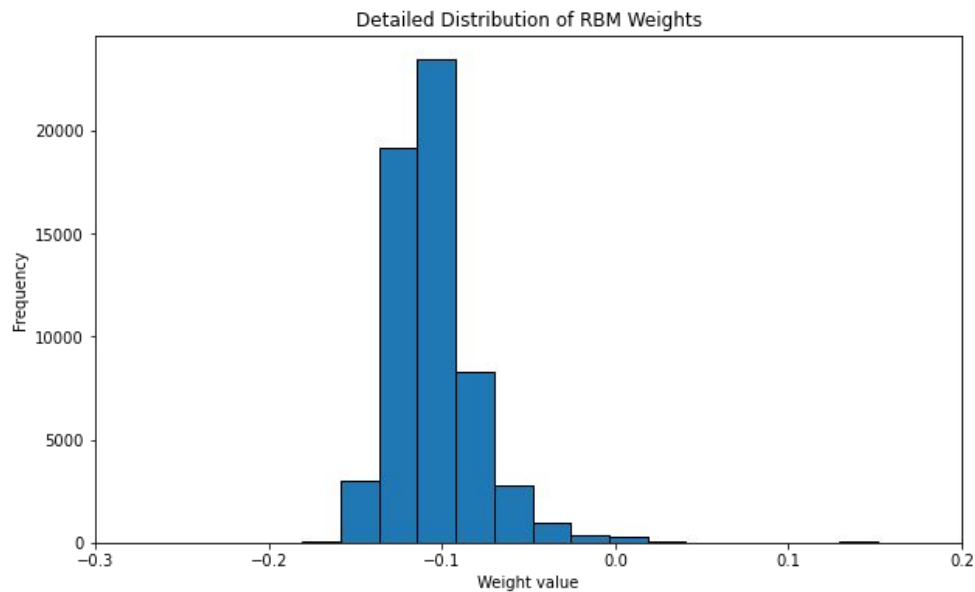
- Key Properties:
  - Gradients must respect the hyperbolic geometry (e.g., Riemannian optimization).
  - Optimization is more challenging due to curvature and numerical instability.

# Implementation

Jupyter notebook, nltk wordnet, numpy, sklearn (for BernoulliRBM), matplotlib, networkx (for visualization)



# I ran both models on the 1.2k synset “mammal” tree



# Results

## Euclidean

50 hidden layers, 50 epochs

Chosen Sample Node:  
pug.n.01

Original active synsets:

['canine.n.02', 'carnivore.n.01', 'dog.n.01', 'mammal.n.01', 'placental.n.01', 'pug.n.01']

Reconstructed active synsets:

['ape.n.01', 'equine.n.01', 'even-toed\_ungulate.n.01', 'horse.n.01', 'mammal.n.01', 'odd-toed\_ungulate.n.01', 'placental.n.01']

Reconstruction error (mean absolute difference): 0.0077

Precision=0.286, Recall=0.333, F1=0.308

## Hyperbolic

50 hidden layers, 3 epochs (~60 min)

Chosen Sample Node:  
pug.n.01

Original active synsets:

['canine.n.02', 'carnivore.n.01', 'dog.n.01', 'mammal.n.01', 'placental.n.01', 'pug.n.01']


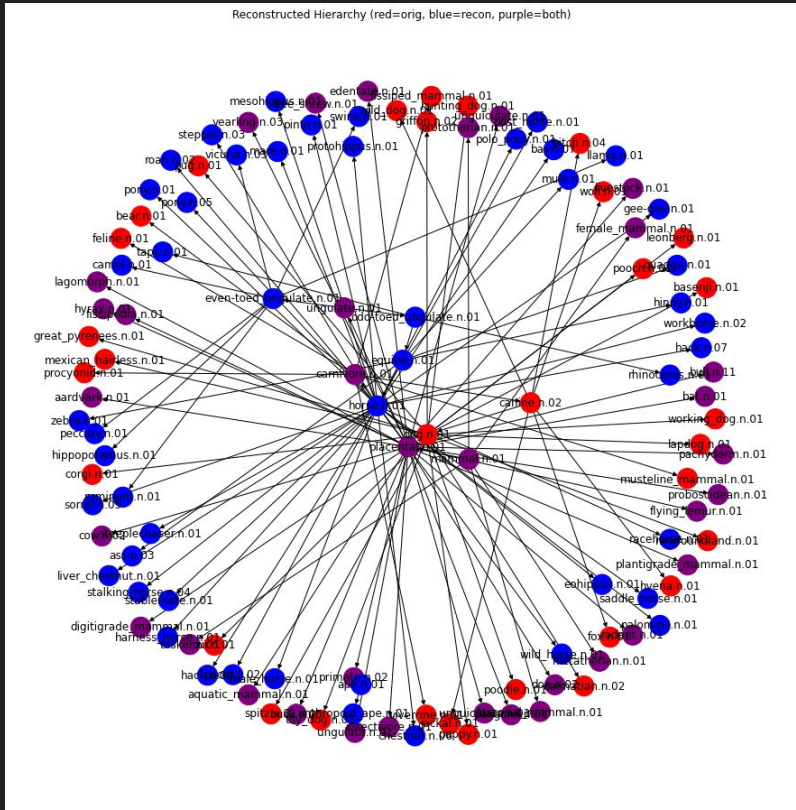
Reconstructed active synsets:

['canine.n.02', 'carnivore.n.01', 'dog.n.01', 'mammal.n.01', 'phyllostomus\_hastatus.n.01', 'placental.n.01']

Reconstruction error (mean absolute difference): 0.0017

Precision=0.833, Recall=0.833, F1=0.833

## Reconstruction Graphs (2 layers deep)



# Results

Now reconstructing 'dog.n.01'

## Euclidean

Chosen Sample Node:

dog.n.01

Original active synsets:

['canine.n.02', 'carnivore.n.01', 'dog.n.01', 'mammal.n.01', 'placental.n.01']

Reconstructed active synsets:

['bovid.n.01', 'dog.n.01', 'mammal.n.01', 'odd-toed\_ungulate.n.01', 'placental.n.01', 'primate.n.02', 'rabbit-eared\_bandicoot.n.01']

Reconstruction error (mean absolute difference): 0.0051

Precision=0.429, Recall=0.600, F1=0.500

## Hyperbolic

Chosen Sample Node:

dog.n.01

Original active synsets:

['canine.n.02', 'carnivore.n.01', 'dog.n.01', 'mammal.n.01', 'placental.n.01']

Reconstructed active synsets:

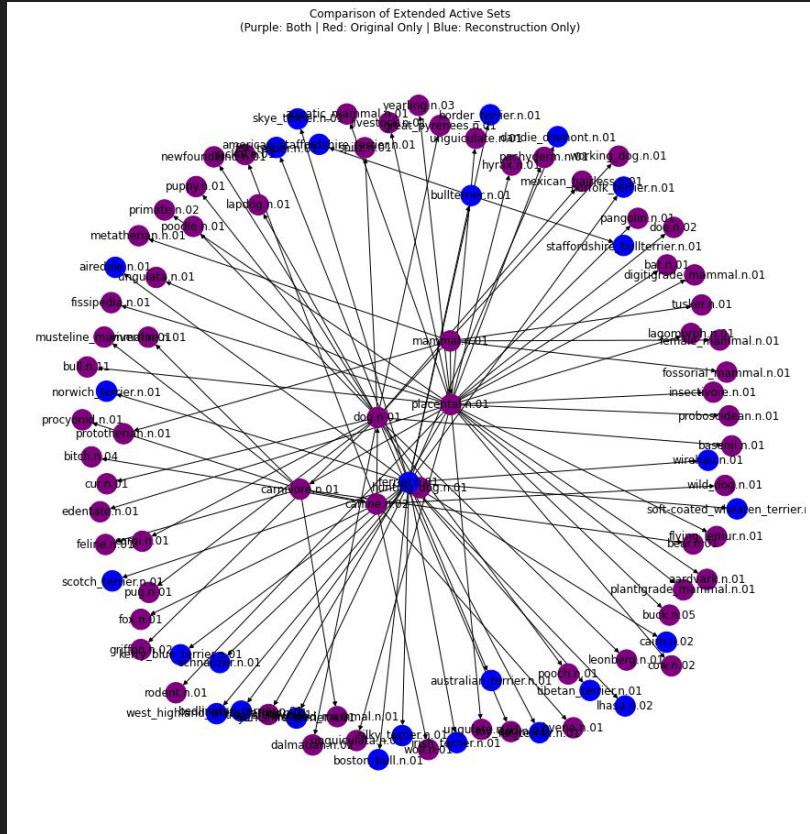

['canine.n.02', 'carnivore.n.01', 'dog.n.01', 'hunting\_dog.n.01', 'mammal.n.01', 'placental.n.01']

Reconstruction error (mean absolute difference): 0.0009

Precision=0.833, Recall=1.000, F1=0.909



## Reconstruction Graphs (2 layers deep)





# Challenges and Future Work

- Training Time
  - Hyperbolic arithmetic is more computationally intensive than standard float operations.
  - Future work: explore optimized implementations or approximations -- rewriting the model using PyTorch would enable GPU acceleration
- Future models could treat curvature as a hyperparameter or optimize it jointly with weights.
- A larger evaluation—e.g., across full WordNet, biomedical ontologies, or hierarchical image taxonomies—would provide stronger evidence of HRBM effectiveness.